

PipedPeer: Transparent Distribution of Unmodified Python Workloads over a Zero-Configuration Peer Network

Yateen Vaviya

Department of Computer Engineering
Sardar Patel Institute of Technology
Mumbai, Maharashtra, India

Vinayak Yadav

Department of Computer Engineering
Sardar Patel Institute of Technology
Mumbai, Maharashtra, India

Sahil Gupta

Department of Computer Engineering
Sardar Patel Institute of Technology
Mumbai, Maharashtra, India

Abstract—Distributed computing frameworks require a cluster that somebody has already configured, and they require the program to be rewritten against their API. Both costs are paid before any speedup is obtained, which puts the idle machines on a small team’s network out of reach for ordinary work. We present PipedPeer, a system that runs an unmodified Python script on another machine on the network with no cluster configuration and nothing installed on any participant beyond a single binary. Two mechanisms make this possible. First, the script’s environment is built as a content-addressed Nix closure against a pinned package set, so the same script yields an identical store path on every machine and on any day; that path is then used as a cluster-wide cache key, turning environment transfer from a per-job cost into a once-per-cluster cost. On a three-node cluster the first submission of a job took 12.00 s and subsequent submissions of the same environment took 0.66 s, an eighteenfold reduction. Second, the parallel primitives the script already uses are substituted at import time by a shim injected into the job’s interpreter, and are distributed only when a cost model fed by a measured bandwidth probe finds that moving the data costs less than computing it locally. The model declines to distribute dense matrix multiplication, and the project treats that refusal as a tested property rather than a missing feature. Every remote path degrades to a working local one: a worker killed during a distributed parallel map left the run to complete with correct results.

Index Terms—distributed computing, peer-to-peer, transparent interception, reproducible environments, content-addressed caching, cost-based scheduling, fault tolerance

I. INTRODUCTION

The gap between running `python script.py` and running the same work on several machines is large in operational effort and often small in the computation involved. Apache Spark [1], Ray [2] and Dask [3] all presuppose a configured cluster with a head node, matching software on every worker, and a scheduler process. Remote execution over `ssh` avoids that setup but provides no environment management, no placement, no isolation and no fault tolerance.

Two distinct obstacles are usually treated separately. The *environment* obstacle is that a remote machine must be made to match the local one, which is the bulk of the work and the reason ad hoc remote execution does not scale past one

machine. The *code* obstacle is that using more than one machine normally requires rewriting the program.

PipedPeer removes both. Its contributions are:

- 1) **The store path as a cluster-wide cache key.** Pinning an exact package set makes an environment’s content-addressed name identical across machines and across time, which is what allows a closure to be transferred once per cluster rather than once per job (Section III-B).
- 2) **Transparent interception under a measured cost model.** Parallel primitives are substituted at import time and distributed only when the arithmetic favours it, with link bandwidth measured rather than assumed (Sections III-C, IV).
- 3) **Exact aggregation without combiners.** Hash-shuffling on the grouping key places every key complete on one node, so any aggregation is exact, including those with no combiner form (Section III-D).
- 4) **Placement without a scheduler.** Hard filters followed by a bounded score, evaluated in the submitting client, with a three-phase lease providing atomic reservation and recovery through expiry (Section III-A).
- 5) **An unconditional degradation guarantee,** enforced structurally rather than by error handling (Section V).

II. RELATED WORK

Spark [1], Ray [2] and Dask [3] are mature and considerably more capable at scale than the system described here; they differ in requiring a configured cluster and, for Spark and to a lesser degree the others, a program written against their API. BOINC [4] pioneered volunteer computing but is a platform for specific projects, each with its own server infrastructure, rather than a general executor. Nix [5] contributes the content-addressed store on which our caching argument rests.

Work on offloading in edge computing addresses the same decision our cost model makes. Deng et al. [6] formulate offloading as a comparison of transmission against local computation time; Yong et al. [7] show partial offloading to beat all-or-nothing strategies, which is the regime our interception operates in. Learned policies [8] exceed what a deterministic rule achieves, at the cost of a training phase




Fig. 1: System design. The submitting side builds the environment and places the job; the worker admits it, caches the closure and runs it sandboxed. Both roles are the same binary.

that a zero-configuration tool cannot assume. Zhao et al. [9] show topology-aware placement to outperform scalar scoring, a limitation our placement function shares and which we identify as future work.

On failure transparency, Bergström [10] surveys abstractions that hide distributed failure from application code; since a user script here contains no failure handling at all, this is the property Section V must provide. The Python Array API [11] shares our goal of portability without code change but inverts the mechanism, asking libraries to converge on one interface rather than substituting the implementations behind interfaces users already call.

III. SYSTEM DESIGN

Every machine runs one daemon; the command-line client is short-lived. The coordinator is a library linked into that client rather than a service, so there is no scheduler process, no head node and no master. A machine is a submitter and a worker simultaneously, and a new participant joins by running the binary.

A. Placement

Candidates are gathered from a local store of known peers, from UDP probes on the subnet, and optionally from a registry. Five hard filters run in order: architecture must match the closure, a GPU must exist if required, and free memory, cores and VRAM must each satisfy the request. Survivors are scored:

$$S = \left(1 - \frac{u_{\text{cpu}}}{200} - \frac{u_{\text{mem}}}{200} - 0.05 n_{\text{jobs}} + 0.10(\hat{c} - \frac{1}{2}) + 0.06(\hat{f} - \frac{1}{2}) + 0.10(\hat{m} - \frac{1}{2}) + 0.10(\hat{r} - \frac{1}{2}) \right) \cdot H \quad (1)$$

where u are utilisations in percent, n_{jobs} is the running job count, $\hat{c}, \hat{f}, \hat{m}, \hat{r}$ are core count, clock, free memory and free-core ratio normalised onto $[0, 1]$, and H is a health score. Two properties matter more than the coefficients. Load dominates hardware: the load terms can subtract 1.0 while the hardware terms move the score by at most 0.36, so a fast but busy machine loses to a slower idle one. And a missing metric normalises to 0.5, the neutral midpoint, so uneven telemetry across heterogeneous machines neither rewards nor penalises a node.

Work is reserved by a three-phase lease. The admission check and the lease insertion occur under one lock, so concurrent submitters cannot both claim the last slot, and the concurrency limit counts reserved and running leases together so a burst cannot slip through the window between the phases. Recovery needs no further mechanism: a submitter that dies stops renewing and its slot is reclaimed, while a worker that dies drops out of the healthy set and its job is rescheduled.

B. Environment Reproduction

The client scans the script's imports, resolves them to packages, generates a Nix flake and builds a closure. The flake pins an *exact* package-set revision rather than a branch, and this is the load-bearing decision of the entire design. A branch resolves at build time, so two machines building the same script on different days would produce different store paths; because the cache is keyed on the store path, the system would re-ship a large archive for every job. An exact pin makes the store path a stable, content-derived name for an environment across every machine and across time.

Before uploading, the client asks the target whether it already holds that path. Two distinct predicates are needed and both are implemented: whether the archive is cached, and whether the closure is materialised on disk. A deployment with a shared store volume satisfies the second without the first, and a per-node store satisfies the first before the second; either predicate alone is wrong for one of the two layouts. After an upload the worker pushes the closure to peers that lack it, which is what makes those peers eligible to receive distributed work.

C. Transparent Interception

A Python module embedded in the binary is appended to the job’s workspace archive and placed on the interpreter’s module path, so it is imported automatically before the user’s first line. Nothing is installed on any node and the user’s project on disk is never modified. Each library is patched on its first import rather than up front, since importing a large framework such as torch costs on the order of a second and charging that to jobs that never use it would violate the guarantee of Section V.

Substituted primitives are parallel maps, pandas grouping, joining and chunked reading, NumPy matrix multiplication and decompositions, joblib and scikit-learn parallelism, and data-parallel training.

D. Exact Aggregation Without Combiners

For grouping and joining, rows are bucketed by a hash of the key. Equal keys hash equally, so every key lands complete on exactly one node. Each node’s aggregation is therefore a complete aggregation over the keys it holds, and the origin combines partials by concatenation alone. No combiner algebra is involved, and consequently *any* aggregation specification is exact, including user-defined reductions that admit no combiner. This is a stronger guarantee than the partial-aggregation strategies used by systems that push combiners to the workers.

E. Data-Parallel Training Without a Socket Mesh

Standard data-parallel training establishes direct connections between every pair of ranks on ports of its own choosing, an assumption that holds in a data centre and often nowhere else. Instead, each rank posts its gradient to the lead node’s ordinary daemon port and receives the full set once every rank has arrived; the daemon never interprets the payload and averaging happens in the shim. This trades bandwidth for requiring nothing beyond the port the rest of the system already uses.

IV. THE COST MODEL

Interception is only worthwhile when the cluster wins, so every intercepted call is gated. Writing b for payload bytes, ϕ for arithmetic intensity in operations per byte, B for measured bandwidth, R for local throughput and K for node count, work is distributed only when

$$\frac{b\phi}{R} > \frac{b}{B} + \frac{b\phi}{RK} \times 1.3 \quad (2)$$

the factor 1.3 charging 30% overhead against ideal parallel speedup. Three guards precede the arithmetic: at least two nodes, a payload of at least 32 MB, and a carve-out rejecting low-intensity work below 512 MB. Bandwidth is measured by pushing 64 MB through the real dispatch path, cached for five minutes; if the probe fails the answer is to stay local.

A separately calibrated model governs dense linear algebra and produces a result worth emphasising: **matrix multiplication is never distributed**. Local BLAS sustains roughly 200 GFLOP/s, which no cluster transport beats at practical




Fig. 2: Decision regions of Equation 2 at three nodes on a 500 MB/s link.

sizes. The project encodes this as a tested property, failing its demonstration suite if a matrix multiplication is ever observed leaving a node. Singular value decomposition, two orders of magnitude slower locally, does move.

V. DEGRADATION GUARANTEE

Using the cluster must never be worse than not using it. Five independent layers enforce this. Chunks of a parallel map are dealt alternately to the local side and the cluster and fill a result table first-wins, with idle local cores re-running any chunk the cluster has not returned; the call returns only once the local side has covered every item, so a wholly unresponsive cluster degrades to a plain local pool. Beneath that, every intercepted call falls back to the original implementation on any exception, a failed peer is skipped in favour of the next, a dead worker process is respawned, and a failed job is rescheduled with exponential backoff.

VI. EVALUATION

A. Method and its limits

Measurements were taken on one host with 20 logical cores running three worker daemons in containers. **This bed cannot measure multi-machine speedup**, because moving processor-bound work between containers on one host cannot make that host faster, and no speedup figure is claimed. It measures correctness, overhead, cache behaviour and fault tolerance, which are the properties the design argument rests on.

B. Environment cache

The same job submitted five times took 12.00 s on the first submission and a median of 0.66 s thereafter, an **eighteenfold reduction** (Fig. 3). The first submission ships the closure; the rest find it materialised and ship only the project files. This is the direct consequence of pinning the package set so that a store path names an environment identically everywhere.




Fig. 3: Cold and warm submission of the same environment.

TABLE I: PipedPeer against existing systems

Property	Ours	Spark	Dask	ssh
Runs unmodified Python	yes	no	partial	yes
No cluster to configure	yes	no	no	partial
No head or master node	yes	no	no	yes
Reproduces the environment	yes	partial	partial	no
Sandboxes the job	yes	no	no	no
Reschedules on node loss	yes	yes	yes	no

C. Cost of leaving interception on

Since interception is on by default, its cost when the model declines to distribute is what most users pay. On a workload the model declines, a whole remote job took 0.93 s against 0.85 s with interception disabled, an overhead of 9%.

Interpreter startup for a job that imports nothing heavy is a separate matter. Eagerly patching libraries cost $21.60\times$ the unshimmed startup, against a budget of $3\times$ the project sets itself; deferring each library’s patching to its first import brings this to $1.20\times$.

D. Correctness and fault tolerance

The suite passes 246 test cases with none failing. Distribution correctness is established by comparison against local execution: grouped aggregation and all four join types are frame-equal to local pandas, and each rank’s gradient in data-parallel training is exactly the mean across ranks.

Under fault injection, a twenty-item parallel map was distributed across the cluster and a worker was killed mid-computation. The run completed with correct results, the checksum matching the serial computation.

E. Threats to validity

The single-host bed is the principal threat and the reason no speedup is claimed. Container workers share a processor, a memory system and the loopback interface, so both transfer costs and contention differ from a real network. The cache result is the least affected, since it turns on whether an archive is sent at all rather than on how fast it travels; the interception overhead figures are the most affected and should be read as lower bounds on the benefit and upper bounds on the relative overhead.

VII. CONCLUSION

PipedPeer runs unmodified Python on other machines with no cluster to configure, by reproducing environments as content-addressed closures whose store path serves as a cluster-wide cache key, and by substituting the parallel primitives a script already uses under a cost model driven by measured bandwidth. The measurements that carry the argument are an eighteenfold reduction in repeat submission cost and correct completion of a distributed computation after a worker was killed mid-run.

The part we consider most transferable is the discipline of the cost model. A system that distributed everything it could intercept would be slower than plain Python on most real workloads, because a modern BLAS or a warm page cache usually beats the network. Measuring rather than assuming, declining to distribute when local wins, and treating that refusal as a property to be tested is what makes always-on interception defensible.

Future work is led by authentication and an encrypted transport, without which the system is confined to a trusted network; then discovery beyond the subnet, locality- and topology-aware placement, and measurement across separate machines to obtain the speedup figures this evaluation deliberately does not claim.

REFERENCES

- [1] M. Zaharia, M. Chowdhury, M. J. Franklin, S. Shenker, and I. Stoica, “Spark: Cluster computing with working sets,” in *Proc. 2nd USENIX Workshop on Hot Topics in Cloud Computing*, 2010.
- [2] P. Moritz et al., “Ray: A distributed framework for emerging AI applications,” in *Proc. 13th USENIX OSDI*, 2018, pp. 561–577.
- [3] M. Rocklin, “Dask: Parallel computation with blocked algorithms and task scheduling,” in *Proc. 14th Python in Science Conf.*, 2015, pp. 130–136.
- [4] D. P. Anderson, “BOINC: A system for public-resource computing and storage,” in *Proc. 5th IEEE/ACM Workshop on Grid Computing*, 2004, pp. 4–10.
- [5] E. Dolstra, M. de Jonge, and E. Visser, “Nix: A safe and policy-free system for software deployment,” in *Proc. 18th LISA*, 2004, pp. 79–92.
- [6] Y. Deng, Z. Chen, X. Chen, and Y. Fang, “Task offloading in multi-hop relay-aided multi-access edge computing,” *IEEE Trans. Veh. Technol.*, vol. 72, no. 1, pp. 1372–1376, 2023.
- [7] D. Yong, R. Liu, X. Jia, and Y. Gu, “Joint optimization of multi-user partial offloading strategy and resource allocation in D2D-enabled MEC,” *Sensors*, vol. 23, no. 5, 2023.
- [8] H. He, X. Yang, X. Mi, H. Shen, and X. Liao, “Multi-agent deep reinforcement learning based dynamic task offloading in a D2D mobile-edge computing network,” *Sensors*, vol. 24, no. 16, 2024.
- [9] Z. Zhao, J. Perazzone, G. Verma, and S. Segarra, “Congestion-aware distributed task offloading in wireless multi-hop networks using graph neural networks,” arXiv:2312.02471, 2023.
- [10] A. Bergström, “Programming models for failure-transparent distributed systems,” M.S. thesis, KTH Royal Institute of Technology, 2025.
- [11] A. Meurer et al., “Python array API standard: Toward array interoperability in the scientific Python ecosystem,” in *Proc. SciPy*, 2023.